


HOW THE DAILY BRIEF ACTUALLY RUNS

AI Concurrency: Fan-Out / Fan-In

One prompt splits into several independent research tasks that run in parallel, then the results funnel back into a single composed answer — written to a file and pushed to the web.

THE PATTERN



Fan-out: five independent tasks launch in one batch. Fan-in: results are gathered once, composed in order, and written out.

THE TASK, BY THE NUMBERS

9

Web searches, issued concurrently

1

Page fetch

~4

Bash calls

8–10

Back-and-forth turns

Plus reading and writing the final HTML brief — all wrapped into one run.

WHAT IT COSTS (JULY 2026 TOKEN PRICING — NOT SUBSCRIPTION PRICING)

Claude Sonnet 5 \$2 / MTok in · \$10 / MTok out (through Aug 31, 2026 — rising to \$3 / \$15)		Claude Haiku 4.5 \$1 / MTok in · \$5 / MTok out	
Input (cumulative, ~125K tok)	≈ \$0.25	Input (cumulative, ~125K tok)	≈ \$0.125
Output (~6K tok)	≈ \$0.06	Output (~6K tok)	≈ \$0.03
Total run		Total run	
\$0.30 – \$0.35		\$0.15 – \$0.18	

Input tokens include the system prompt and tool schemas resent every turn — the biggest cost driver in a multi-turn agentic run.

THE TAKEAWAY

Fan-out / fan-in isn't just faster — it's **the same total work, done in parallel instead of in sequence**. Instead of nine searches waiting on each other one at a time, they all run at once, and the model composes the final answer only after everything comes back. On a subscription, the marginal cost of a run like this is a fraction of a cent to a few tenths of a cent — and it's genuinely fun to watch all the tasks fire off together.